

Apache Spark

AmirReza Mohammadi

2019 winter


Apache Spark



- *Lightning-fast unified analytics engine* for large-scale data processing.

- Speed:

Run workloads 100x faster.




Apache Spark



- *Lightning-fast unified analytics engine* for large-scale data processing.
- Speed:
 - Run workloads 100x faster.
- Ease of Use
 - Write applications quickly in Java, Scala, Python, R, and SQL.
- Generality
 - Combine SQL, streaming, and complex analytics.

Spark Components



Spark.apache.org

Spark SQL

Apache Spark's module for working with structured data.

- Integrated

Seamlessly mix SQL queries with Spark programs.

```
results = spark.sql(  
    "SELECT * FROM people")  
names = results.map(lambda p: p.name)
```

Apply functions to results of SQL queries.

Spark.apache.org

Spark SQL

Apache Spark's module for working with structured data.

- Uniform Data Access

Connect to any data source the same way.

```
spark.read.json("s3n://...")  
  .registerTempTable("json")  
results = spark.sql(  
    """SELECT *  
       FROM people  
       JOIN json ...""")
```

Query and join different data sources.
Spark.apache.org

Spark Streaming

Spark Streaming makes it easy to build scalable fault-tolerant streaming applications.

- Ease of Use

Build applications through high-level operators.

```
TwitterUtils.createStream(...)  
    .filter(_.getText.contains("spark"))  
    .countByWindow(Seconds(5))
```

Counting tweets on a sliding window

Spark.apache.org

Spark Streaming

Spark Streaming makes it easy to build scalable fault-tolerant streaming applications.

- Spark Integration

Combine streaming with batch and interactive queries.

```
stream.join(historicCounts).filter {  
  case (word, (curCount, oldCount)) =>  
    curCount > oldCount  
}
```

Find words with higher frequency than
historic data

Spark MLlib

- **MLlib** is Apache Spark's scalable machine learning library.
- Ease of Use

Usable in Java, Scala, Python, and R.

```
data = spark.read.format("libsvm")\
    .load("hdfs://...")
```

```
model = KMeans(k=10).fit(data)
```

Calling MLlib in Python

Spark.apache.org

Spark MLlib

ML algorithms include:

- Classification: logistic regression, naive Bayes,...
- Regression: generalized linear regression, survival regression,...
- Decision trees, random forests, and gradient-boosted trees
- Recommendation: alternating least squares (ALS)
- Clustering: K-means, Gaussian mixtures (GMMs),...
- Topic modeling: latent Dirichlet allocation (LDA)
- Frequent itemsets, association rules, and sequential pattern mining

Spark MLlib

ML workflow utilities include:

- Feature transformations: standardization, normalization, hashing,...
- ML Pipeline construction
- Model evaluation and hyper-parameter tuning
- ML persistence: saving and loading models and Pipelines

Other utilities include:

- Distributed linear algebra: SVD, PCA,...
- Statistics: summary statistics, hypothesis testing,...

Spark GraphX

GraphX is Apache Spark's API for graphs and graph-parallel computation.

- Flexibility

Seamlessly work with both graphs and collections.

- Algorithms

Choose from a growing library of graph algorithms.

- Speed

Comparable performance to the fastest specialized graph processing systems.